

MACHINE LEARNING - U3.

Q. Explain Statistical Learning Theory

Ans. Statistical Learning Theory (SLT) is a framework that provides the mathematical function for machine learning. It uses concepts from statistics and functional analysis to explain how a learning algorithm can make valid predictions based on a finite sets of data.

The goal is to understand how well the model is trained on past data to generalize new and unseen data.

The Central Problem

Imagine you want to build a model that can predict a home's price based on its size. You collect data on the square footage and selling price of many homes. A learning algorithm uses this data to find a function (model) that best describes the relationship between the size and price.

The central problem is that this function ^{must} works well not just for the homes you used to train the model, but also for all future, unseen home (the test data). SLT provides the tools to analyze and guarantee this ability to generalize.

The Bias Variance Trade-off, this is one of the most important concept in SLT and machine learning. It describes the conflict between the two types of errors that a model can make.

① BIAS:

The error from a model that is too simplistic. A simple model might ~~be~~ underfit the training data, essentially and fail to capture important patterns.

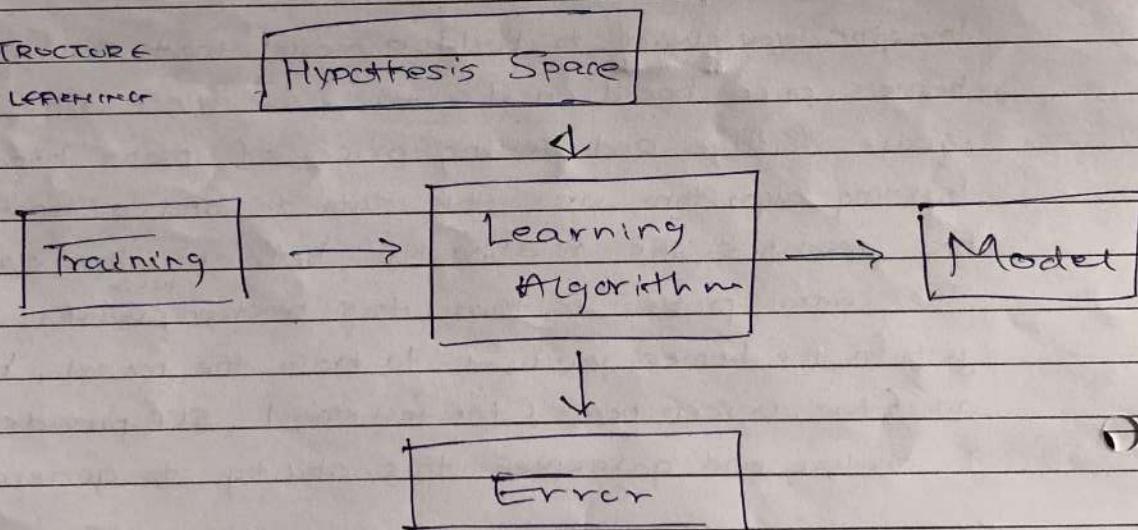
Ex: Trying to fit a straight line into complex curved relationship.

① VARIANCE

The error from a model that is too complex, a complex model can overfit the training data, essentially memorizing the noise and random fluctuations rather than the true underlying relationship. This makes the model perform poorly on new data.

The goal is to find the right balance : A model that is complex enough to capture the true pattern but not so complex that it also memorizes the noise.

• STRUCTURE OF LEARNING



② Hypothesis Space

A collection of all possible models (e.g. all possible linear models for or all possible neural models for) that the learning algorithm can choose from

③ Learning Algorithm

A procedure to select the "best" model from the hypothesis space based on the training data

⑩ Error / Risk

A metric used to evaluate how well a model performs.

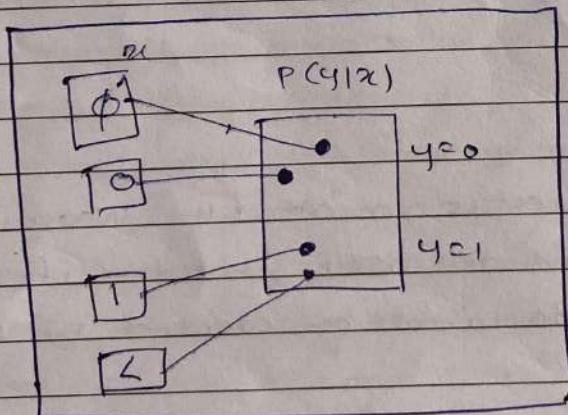
SLT distinguishes between two types

TRAINING ERROR & TEST ERROR.

This Theory provides guarantee that for a "well-behaved" learning problem, the training error would be approx. of test error.

⑪ Generative Model & It's Formulation

FNS A generative model is a type of machine learning model that aims to learn underlying patterns or distribution of data to generate new similar kind of data. This is used in unsupervised machine learning algorithm to describe phenomena in data, enabling computers to understand real world.



Ex: Imagine you're teaching a child to draw animals. After showing them several pictures of different animals, the child begins to understand the general features of each animal. Given some time, the child might draw an animal they've never seen before combining features they've learned. This is

analogous to how a generative model operates: It learns from the data it is been exposed to and then creates something new based on that knowledge.

The model of distribution of data in statistics can be $P(x)$ or Jointly $P(x,y)$

The Formulation of Generative Model can be described representing using different different models, each representing a different approach to learning the underlying data distribution. $P(x)$.

1. GANs

This method involves of two AI systems competing with each other.

① Generator: This is the artist, or "forger" whose job is to create new image (or text / Audio / Video) that looks as real as possible

② The Discriminator

This is the detective, or "art crit", whose job is to look at a picture and tell itself, whether if it's a real one from the training data or fake one created by Generator

These Two AI's train together in the game of one-upmanship

- Generator Creates a Fake Image

- Discriminator Tries to catch it

Both model learns and improve from the process. The Generator learns to create fake and Discriminator learns to become better at detecting them.

(Q) Bayesian Decision Theory

Ans. Bayesian Decision Theory is an statistical approach to used to make the best decision theory under uncertainty. It combines probability to represent uncertain loss function to measure the cost of evidence. The decision rule is to choose the action with the lowest expect loss. Bayes' Theorem, it updates beliefs about different outcomes based on available evidence. It provides an optimal and logical framework for decision-making in ML

This process begins with three main components: concepts.

(i) Prior probability (ii) Likelihood (iii) posterior probability.

The prior probability $p(w)$ represent what we already believe about an event or class without seeing any new data.
Ex: In medical Diagnosis, it could be the overall chance that a person has disease in the general population.

The likelihood $P(x|w)$ tells us how likely it is to observe the data x if the true class is w .

Ex: This means how likely a person's test result is if they actually have disease.

Finally, the posterior probability $p(w|x)$ is the updated belief about the event after seeing the data, and can be calculated as

$$p(w|x) = \frac{p(x|w) p(w)}{p(x)}$$

Eventually, the generator becomes so good that the discriminator can no longer tell the fakes from the real thing.

At this point, the Generator has learned how to create very realistic own content

3. Variational Interfaces - VAE's

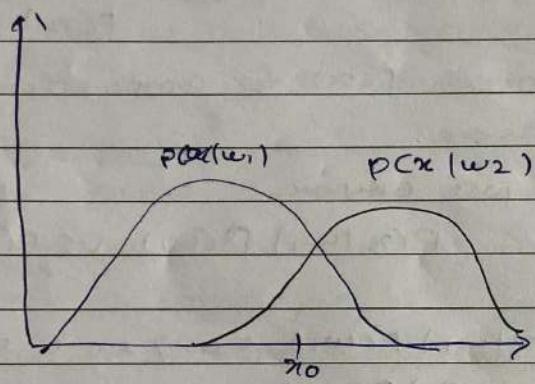
This method introduces a probabilistic twist to autoencoder architecture. It assumes that data, " x " is generated from some underlying, simpler latent variable, " z ", via a decoder network " $p(x|z)$ ". Since the true latent distribution $p(z|x)$ is unknown, a VAE uses an encoder network $q_\phi(z|x)$ to approximate it.

Q. Bayesian Decision Theory.

Ans. Bayesian theorem is a fundamental in machine learning, especially in the context of Bayesian Inference. It provides a way to update our beliefs about hypothesis on new evidence.

Bayes' Theorem is a fundamental concept of probability theory, that plays an crucial role in various ML algorithms, especially, in the field Bayesian statistics and probabilistic modelling. It provides an way to update probabilities based on new evidence or information. In context of ML, Bayes' theorem is often used in probabilistic model.

This theorem updates our prior belief based on new evidence from the data, once we know the posterior probab the Bayesian Decision Rule tells us to choose the class with highest posterior probab, meaning the class that is most likely the evidence. Alternatively if the loss or costs are defined, we choose the action that results in smallest expected loss



Graphically, Bayesian Decision Theory can be illustrated using two probability curves. Each curve shows how likely a value of x is for each class. The point where these curves intersect is called the decision boundary (x_0). If a data point lies on (x_0) this mean both classes are likely. If data point lies on one side of the boundary it is classified as class 1 and on other side as class 2. The area under each curve represents the probab of diff outcome for each class.

Beauty of this Theorem is it ensures optimal Decision under uncertainty. Even when outcomes are not guaranteed this model provides the most logical and risk balanced choice based on all available info ...

Ex: We want to classify a sample x into one of 2 classes

- Class w_1 (e.g. "spam mail")
- Class w_2 (e.g. "not spam")

Given

$$P(w_1) = 0.4 \quad P(w_2) = 0.6$$

$$P(x|w_1) = 0.7 \quad P(x|w_2) = 0.2$$

$$P(w_1|x) = \frac{P(x|w_1) P(w_1)}{P(x)}$$

We don't need $P(x)$ for comparison, since it's same for both classes.

So we will just compare

$$P(x|w_1) P(w_1) \text{ vs } P(x|w_2) P(w_2)$$

$$P(x|w_1) P(w_1) = 0.7 \times 0.4 = 0.28$$

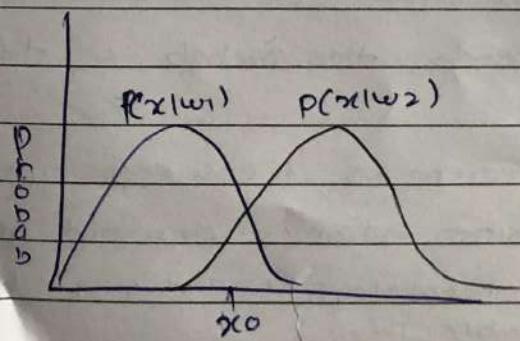
$$P(x|w_2) P(w_2) = 0.2 \times 0.6 = 0.12$$

Decision Rule

Choose the class with higher posterior product

$$0.28 > 0.12 \Rightarrow \text{Decide } w_1$$

So, the feature x is classified as belonging to class w_1 = Spam mail



Q. Gaussian Mixture Model

AHS.

Gaussian Mixture Model is a powerful statistical method used to represent data that seems to come from multiple subgroups or patterns. Instead of assuming that all data points come from one single source / Gaussian ~~model~~ (Normal) Distribution, GMM assumes that the data is generated from a mixture of several Gaussian Distribution, each representing different cluster or group in the dataset. This makes it an excellent tool for modelling complex, real-world data that cannot be described accurately by one bell-shaped curve.

Each Gaussian model component is an mixture defined by three parameter: a mean (μ) that represents the centre of the cluster, a covariance (Σ) that measures the shape of cluster, and a mixing coefficient (π) that tells how much that Gaussian model contributes to the overall distribution. Together these parameters define how mixtures behaves. Mathematically, the probability density function for a Gaussian mixture model is written as

$$p(x^n) = \sum_{k=1}^K \pi_k n(\mu_k, \Sigma_k)$$

Here K is the number of Gaussian Components. π_k is the weight (or prior prob) of the k^{th} component. $n(\mu_k, \Sigma_k)$ is the Gaussian Distribution for that Component. The sum of all weights equals 1, meaning the model considers all component together as a part of whole.

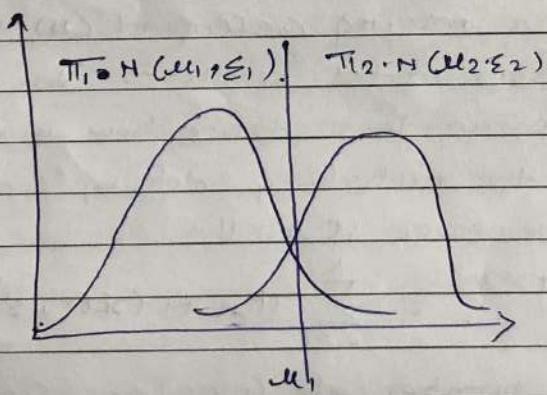
The goal is to estimate these parameters so that the overall model best fits the observed data. This estimation is typically done using the Expectation-Maximization (EM) algorithm, an iterative process that alternates between two steps:

(1) EXPECTATION (E-STEP)

Calculates the probability that each data point belongs to each Gaussian component based on the current parameters.

(2) MAXIMIZATION (M-STEP)

Updates the parameters of each Gaussian distribution and the mixing coefficient to maximize the likelihood of the data given these probabilities.



By repeating these steps until convergence, the GMM learns a set of Gaussian distributions that collectively describes the data's structure.

Q. ~~Generative Models in a Nutshell~~

Q. Expectation maximization method

Ans. The Expectation-Maximization (EM) algorithm is an iterative method used in machine learning to find the maximum likelihood estimates of parameters when data contains hidden or missing variables. It alternates between the two steps ① Expectation - E step

② maximization - M step

① E step

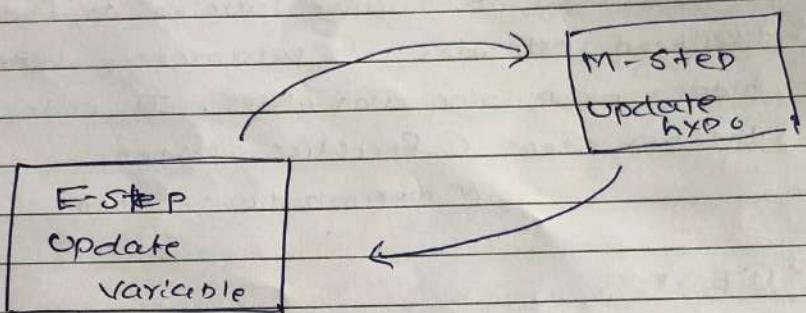
In E-step, the algorithm uses the current parameters estimates to calculate the expected values of the hidden variables. It determines the probability that each data point belongs to each possible hidden state. This means the model estimates how likely each cluster or component is to have generated a particular observation. Essentially, this step fills in missing information probabilistically based on current model

② M step

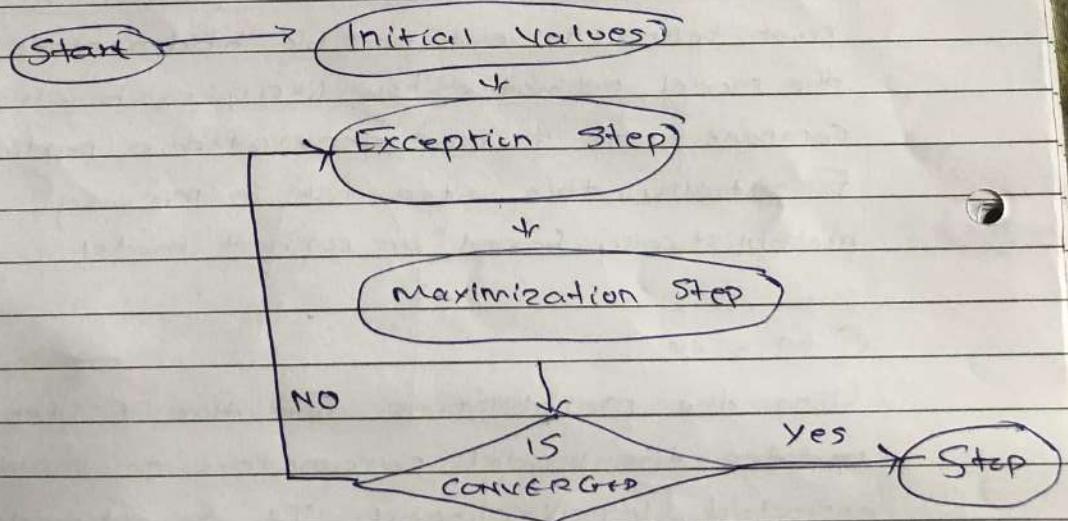
Using the probabilities from the E-step, the M-step updates the model parameters to maximize the expected log-likelihood. It re-estimates parameters like means, variance, or mixing weights so the model better fits the observed data. Each parameter is updated as weighted average, with weights coming from E-Step's probabilities. This step improves model accuracy before next iteration

These two steps are repeated until convergence, which typically means that.

- The parameter value stops changing significantly or
- The log likelihood improves only by a negligible amount



WORKING OF EM



① Initialization

Begin with initial estimates for the model parameters, such as mean, variance and mixing weights. These serve as the starting point for refining the model through iterative updates.

(i) E-Step

Calculate the probability that each data point belongs to each hidden or latent component. This step estimates missing data using the current parameter values

(ii) M-Step

Updates the model parameters to maximize the expected log likelihood from the E-step Results

(iii) Convergence

Check if changes in parameters or log likelihood are below a small threshold. If yes the algorithm stops otherwise it repeats E & M steps